Coding

7.2. Convolutions for Images

7.2.1. The Cross-Correlation Operation

```
In [2]: import torch
        from torch import nn
        from d2l import torch as d2l
In [3]: def corr2d(X, K): #@save
            h, w = K.shape
            Y = torch.zeros((X.shape[0] - h + 1, X.shape[1] - w + 1))
            for i in range(Y.shape[0]):
                for j in range(Y.shape[1]):
                    Y[i, j] = (X[i:i + h, j:j + w] * K).sum()
            return Y
In [4]: X = torch.tensor([[0.0, 1.0, 2.0],
                          [3.0, 4.0, 5.0],
                          [6.0, 7.0, 8.0]])
        K = torch.tensor([[0.0, 1.0], [2.0, 3.0]])
        corr2d(X, K)
Out[4]: tensor([[19., 25.],
                 [37., 43.]])
```

• I understood that operation which is similar to scalar multiplication of Matrix...

7.2.2. Convolutional Layers

```
In [5]: class Conv2D(nn.Module):
    def __init__(self, kernel_size):
        super().__init__()
        self.weight = nn.Parameter(torch.rand(kernel_size))
        self.bias = nn.Parameter(torch.zeros(1))

def forward(self, x):
    return corr2d(x, self.weight) + self.bias
```

7.2.3. Object Edge Detection in Images

```
In [6]: X = torch.ones((6, 8))
       X[:, 2:6] = 0
Out[6]: tensor([[1., 1., 0., 0., 0., 0., 1., 1.],
                [1., 1., 0., 0., 0., 0., 1., 1.],
                [1., 1., 0., 0., 0., 0., 1., 1.],
                [1., 1., 0., 0., 0., 0., 1., 1.],
                [1., 1., 0., 0., 0., 0., 1., 1.],
                [1., 1., 0., 0., 0., 0., 1., 1.]]
In [7]: K = torch.tensor([[1.0, -1.0]])
In [8]: Y = corr2d(X, K)
Out[8]: tensor([[ 0., 1., 0., 0., 0., -1., 0.],
                [0., 1., 0., 0., 0., -1., 0.],
                [0., 1., 0., 0., 0., -1., 0.],
                [0., 1., 0., 0., 0., -1., 0.],
                [0., 1., 0., 0., 0., -1., 0.],
                [0., 1., 0., 0., 0., -1., 0.]
In [9]: corr2d(X.t(), K)
```

7.2.4. Learning a Kernel

```
In [10]: conv2d = nn.LazyConv2d(1, kernel_size=(1, 2), bias=False)
         X = X.reshape((1, 1, 6, 8))
         Y = Y.reshape((1, 1, 6, 7))
         lr = 3e-2
         for i in range(10):
             Y_hat = conv2d(X)
             l = (Y_hat - Y) ** 2
             conv2d.zero_grad()
             l.sum().backward()
             conv2d.weight.data[:] -= lr * conv2d.weight.grad
             if (i + 1) % 2 == 0:
                 print(f'epoch {i + 1}, loss {l.sum():.3f}')
        epoch 2, loss 2.008
        epoch 4, loss 0.497
        epoch 6, loss 0.149
        epoch 8, loss 0.052
        epoch 10, loss 0.020
In [11]: conv2d.weight.data.reshape((1, 2))
Out[11]: tensor([[ 0.9799, -1.0081]])
```

7.3. Padding and Stride

7.3.1. Padding

• If the width and height of the kernel are different, we can opt for different number of paddings for the width and height!

7.3.2. Stride

```
In [15]: conv2d = nn.LazyConv2d(1, kernel_size=3, padding=1, stride=2)
comp_conv2d(conv2d, X).shape

Out[15]: torch.Size([4, 4])

In [16]: conv2d = nn.LazyConv2d(1, kernel_size=(3, 5), padding=(0, 1), stride=(3, 4))
comp_conv2d(conv2d, X).shape

Out[16]: torch.Size([2, 2])
```

• We can stride to even two dimensions!

7.4. Multiple Input and Multiple Output Channels

7.4.1. Multiple Input Channels

```
In [17]: import torch
from d2l import torch as d2l

In [18]: def corr2d_multi_in(X, K):
    return sum(d2l.corr2d(x, k) for x, k in zip(X, K))

In [19]: X = torch.tensor(
        [[[0.0, 1.0, 2.0], [3.0, 4.0, 5.0], [6.0, 7.0, 8.0]],
        [[1.0, 2.0, 3.0], [4.0, 5.0, 6.0], [7.0, 8.0, 9.0]]])

K = torch.tensor(
        [[[0.0, 1.0], [2.0, 3.0]],
        [[1.0, 2.0], [3.0, 4.0]]])

corr2d_multi_in(X, K)
Out[19]: tensor([[ 56., 72.],
        [[104., 120.]])
```

7.4.2. Multiple Output Channels

```
In [20]: def corr2d_multi_in_out(X, K):
    return torch.stack([corr2d_multi_in(X, k) for k in K], 0)

In [21]: K = torch.stack((K, K+1, K+2), 0)
    K.shape

Out[21]: torch.Size([3, 2, 2, 2])
```

• We can make three output channels by concatenateing K, K+1, K+2 kernel tensors!

7.4.3. 1 x 1 Convolutional Layer

7.5. Pooling

```
import torch
from torch import nn
from d2l import torch as d2l
```

7.5.1. Maximum Pooling and Average Pooling

```
In [26]: def pool2d(X, pool_size, mode='max'):
             ph, pw = pool size
             Y = torch.zeros((X.shape[0] - p_h + 1, X.shape[1] - p_w + 1))
             for i in range(Y.shape[0]):
                 for j in range(Y.shape[1]):
                     if mode == 'max':
                         Y[i, j] = X[i: i+p_h, j: j+p_w].max()
                     elif mode == 'avg':
                         Y[i, j] = X[i: i+p_h, j: j+p_w].mean()
             return Y
In [27]: X = torch.tensor([[0.0, 1.0, 2.0],
                           [3.0, 4.0, 5.0],
                           [6.0, 7.0, 8.0]])
         pool2d(X, (2, 2))
Out[27]: tensor([[4., 5.],
                 [7., 8.]])
In [28]: pool2d(X, (2, 2), 'avg')
Out[28]: tensor([[2., 3.],
                  [5., 6.]])
         7.5.2. Padding and Stride
In [29]: X = torch.arange(16, dtype=torch.float32).reshape((1, 1, 4, 4))
         Χ
Out[29]: tensor([[[[ 0., 1., 2., 3.],
                   [4., 5., 6., 7.],
                   [8., 9., 10., 11.],
                   [12., 13., 14., 15.]]])
In [30]: pool2d = nn.MaxPool2d(3)
         pool2d(X)
```

```
Out[30]: tensor([[[[10.]]]])
In [31]: pool2d = nn.MaxPool2d(3, padding=1, stride=2)
         pool2d(X)
Out[31]: tensor([[[[ 5., 7.],
                   [13., 15.]]])
In [32]: pool2d = nn.MaxPool2d((2, 3), stride=(2, 3), padding=(0, 1))
         pool2d(X)
Out[32]: tensor([[[[ 5., 7.],
                   [13., 15.]]])
         7.5.3. Multiple Channels
In [33]: X = torch.cat((X, X+1), 1)
Out[33]: tensor([[[[ 0., 1., 2., 3.],
                   [4., 5., 6., 7.],
                   [8., 9., 10., 11.],
                   [12., 13., 14., 15.]],
                  [[ 1., 2., 3., 4.],
                   [5., 6., 7., 8.],
                   [ 9., 10., 11., 12.],
                   [13., 14., 15., 16.]]])
          • The number of output channels are same before!
         pool2d = nn.MaxPool2d(3, padding=1, stride=2)
In [34]:
         pool2d(X)
```

```
Out[34]: tensor([[[[ 5., 7.], [13., 15.]], [[ 6., 8.], [14., 16.]]]])
```

7.6. Convolutional Neural Networks (LeNet)

```
In [35]: import torch
from torch import nn
from d2l import torch as d2l
```

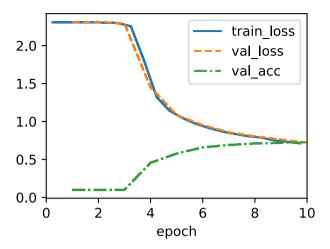
7.6.1. LeNet

In [37]: @d2l.add_to_class(d2l.Classifier) #@save
 def layer_summary(self, X_shape):
 X = torch.randn(*X_shape)

```
In [36]: def init_cnn(module): #@save
             if type(module) == nn.Linear or type(module) == nn.Conv2d:
                 nn.init.xavier uniform (module.weight)
         class LeNet(d2l.Classifier): #@save
             def __init__(self, lr=0.1, num_classes=10):
                 super(). init ()
                 self.save_hyperparameters()
                 self.net = nn.Sequential(
                     nn.LazyConv2d(6, kernel_size=5, padding=2),
                     nn.Sigmoid(),
                     nn.AvgPool2d(kernel_size=2, stride=2),
                     nn.LazyConv2d(16, kernel_size=5),
                     nn.Sigmoid(),
                     nn.AvgPool2d(kernel_size=2, stride=2),
                     nn.Flatten(),
                     nn.LazyLinear(120), nn.Sigmoid(),
                     nn.LazyLinear(84), nn.Sigmoid(),
                     nn.LazyLinear(num classes))
```

```
for layer in self.net:
         X = layer(X)
         print(layer. class . name , 'output shape:\t', X.shape)
 model = LeNet()
model.layer_summary((1, 1, 28, 28))
Conv2d output shape:
                        torch.Size([1, 6, 28, 28])
Sigmoid output shape:
                        torch.Size([1, 6, 28, 28])
AvgPool2d output shape: torch.Size([1, 6, 14, 14])
Conv2d output shape:
                         torch.Size([1, 16, 10, 10])
Sigmoid output shape:
                         torch.Size([1, 16, 10, 10])
AvgPool2d output shape: torch.Size([1, 16, 5, 5])
Flatten output shape:
                         torch.Size([1, 400])
Linear output shape:
                         torch.Size([1, 120])
Sigmoid output shape:
                         torch.Size([1, 120])
                         torch.Size([1, 84])
Linear output shape:
Sigmoid output shape:
                        torch.Size([1, 84])
Linear output shape:
                         torch.Size([1, 10])
```

7.6.2. Training

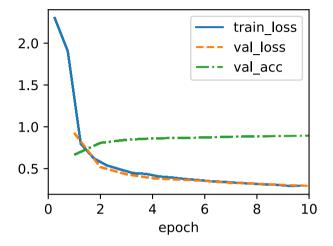


8.2. Networks Using Blocks (VGG)

```
import torch
from torch import nn
from d2l import torch as d2l
```

8.2.1. VGG Blocks

```
super(). init ()
                 self.save_hyperparameters()
                 conv blks = []
                 for (num convs, out_channels) in arch:
                      conv blks.append(vgg block(num convs, out channels))
                  self.net = nn.Sequential(
                     *conv_blks, nn.Flatten(),
                     nn.LazyLinear(4096), nn.ReLU(), nn.Dropout(0.5),
                     nn.LazyLinear(4096), nn.ReLU(), nn.Dropout(0.5),
                     nn.LazyLinear(num classes))
                 self.net.apply(d2l.init cnn)
In [42]: VGG(arch=((1, 64),
                    (1, 128),
                    (2, 256),
                    (2, 512),
                    (2, 512))).layer summary(
              (1, 1, 224, 224))
        Sequential output shape:
                                          torch.Size([1, 64, 112, 112])
        Sequential output shape:
                                         torch.Size([1, 128, 56, 56])
                                         torch.Size([1, 256, 28, 28])
        Sequential output shape:
                                         torch.Size([1, 512, 14, 14])
        Sequential output shape:
        Sequential output shape:
                                          torch.Size([1, 512, 7, 7])
                                  torch.Size([1, 25088])
        Flatten output shape:
        Linear output shape:
                                  torch.Size([1, 4096])
                                  torch.Size([1, 4096])
        ReLU output shape:
        Dropout output shape:
                                  torch.Size([1, 4096])
                                  torch.Size([1, 4096])
        Linear output shape:
        ReLU output shape:
                                  torch.Size([1, 4096])
        Dropout output shape:
                                  torch.Size([1, 4096])
        Linear output shape:
                                  torch.Size([1, 10])
In [43]: model = VGG(arch=((1, 16),
                            (1, 32),
                            (2, 64),
                            (2, 128),
                            (2, 128)),
                      lr=0.01
```



8.6. Residual Networks (ResNet) and ResNeXt

```
import torch
from torch import nn
from torch.nn import functional as F
from d2l import torch as d2l
```

8.6.2. Residual Blocks

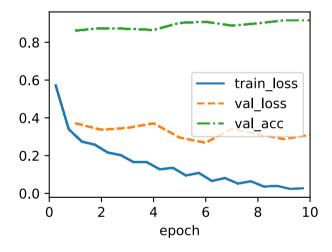
```
kernel size=3,
                                            padding=1)
                 if use 1x1conv:
                     self.conv3 = nn.LazyConv2d(num_channels,
                                                kernel_size=1,
                                                 stride=strides)
                 else:
                     self.conv3 = None
                 self.bn1 = nn.LazyBatchNorm2d()
                 self.bn2 = nn.LazyBatchNorm2d()
             def forward(self, X):
                 Y = F.relu(self.bn1(self.conv1(X)))
                 Y = self.bn2(self.conv2(Y))
                 if self.conv3:
                     X = self.conv3(X)
                 Y += X
                 return F.relu(Y)
In [46]: blk = Residual(3)
         X = torch.randn(4, 3, 6, 6)
         blk(X).shape
Out[46]: torch.Size([4, 3, 6, 6])
In [47]: blk = Residual(6, use_1x1conv=True, strides=2)
         blk(X).shape
Out[47]: torch.Size([4, 6, 3, 3])
         8.6.3. ResNet Model
In [48]: class ResNet(d2l.Classifier):
             def b1(self):
                 return nn.Sequential(
                     nn.LazyConv2d(64, kernel_size=7, stride=2, padding=3),
                     nn.LazyBatchNorm2d(), nn.ReLU(),
```

nn.MaxPool2d(kernel_size=3, stride=2, padding=1))

```
In [49]: @d2l.add_to_class(ResNet)
         def block(self, num residuals, num channels, first block=False):
             blk = []
             for i in range(num residuals):
                 if i == 0 and not first block:
                      blk.append(Residual(num channels,
                                         use_1x1conv=True,
                                          strides=2))
                 else:
                      blk.append(Residual(num channels))
             return nn.Sequential(*blk)
In [50]: @d2l.add_to_class(ResNet)
         def __init__(self, arch, lr=0.1, num_classes=10):
             super(ResNet, self). init ()
             self.save hyperparameters()
             self.net = nn.Sequential(self.b1())
             for i, b in enumerate(arch):
                 self.net.add_module(f'b{i+2}',
                                     self.block(*b, first_block=(i==0)))
             self.net.add_module('last', nn.Sequential(
                 nn.AdaptiveAvgPool2d((1, 1)), nn.Flatten(),
                 nn.LazyLinear(num_classes)))
             self.net.apply(d2l.init cnn)
In [51]: class ResNet18(ResNet):
             def __init__(self, lr=0.1, num_classes=10):
                 super().__init__(((2, 64),
                                   (2, 128),
                                   (2, 256),
                                   (2, 512)),
                                  lr,
                                  num classes)
         ResNet18().layer_summary((1, 1, 96, 96))
```

```
Sequential output shape: torch.Size([1, 64, 24, 24])
Sequential output shape: torch.Size([1, 64, 24, 24])
Sequential output shape: torch.Size([1, 128, 12, 12])
Sequential output shape: torch.Size([1, 256, 6, 6])
Sequential output shape: torch.Size([1, 512, 3, 3])
Sequential output shape: torch.Size([1, 10])
```

8.6.4. Training



Discussions

7.1. From Fully Connected Layers to Convolutions

7.1.1. Invariance

- The CNN takes advantage of translation invariance, which means it is not related to objects' location to find or detect them!
- Then, a first layer of CNN focuses on "local regions" by only referring marginal pixels, which can represent locality!
- Using these factors, we can reduce the number of parameters without constraining the performance.

7.1.4. Channels

- We have to concentrate on the channels because images do not consist of two dimension, but three dimension which contains R, G, B!
- In addition, we can add another channels to the network!

7.2. Convolutions for Images

7.2.2. Convolutional Layers

• Input tensor * Kernel tensor + bias -> Output tensor...

7.2.4. Learning a Kernel

• We can learn the kernel tensor as well by putting inputs and output pairs!

7.2.5. Cross-Correlation and Convolution

• We can get the same outputs of convolution layer regardless of operating cross-correlation or convolution!

7.2.6. Feature Map and Receptive Field

- We can define the output of convolution layer as a feature map!
- Receptive field means that the area of input that contributes to a certain feature map!

7.3. Padding and Stride

7.3.1. Padding

- If we would like to preserve information where is at the edges or corners of image, we can introduce padding information!
- In CNN, we commonly use kernels with odd number height and width (3, 5, 7, ...) to preseve dimension while using same number of rows on top and bottom!

7.3.2. Stride

• To make computational efficiency or downsampling of images, we can skip pixels in the intermediate position by moving 2 or more elements at a time. We can call this method "stride"!

7.4. Multiple Input and Multiple Output Channels

7.4.1. Multiple Input Channels

- So far, we have dealt with the single input and single output channels, where we can consider the input and output tensor as matrix!
- However, if we woule like to take typical images that consist of RGB into account, we should know how to deal with multiple input and multiple output channels!

7.4.3. 1x1 Convolutional Layer

• We can increase nonlinearity by introducing 1x1 convolutional layer!

7.5. Pooling

- Pooling layers can alleviate sensitivity of location and spatially downsampling representations!
- We can easily think about pooling layers as the feature maps of feature map!

7.5.1. Maximum Pooling and Average Pooling

- Average pooling is very similar to image downsampling method. This method takes average value of adjacent pixels instead of choosing the value of second pixel!
- Max pooling is more preferred than average pooling method...

7.5.2. Padding and Stride

• As with convolutional layers, pooling layers can make padding or stride to change output shape!

7.5.3. Multiple Channels

• Unlike convolutional layers, pooling layers pools each input channel seperately, which means the number of output channels of pooling layer and that of input channels!

7.6. Convolutional Neural Networks (LeNet)

- LeNet extracts features through several convolutional layers and pooling layers!
- What if we apply LeNet to large-size images instead of 28 x 28 MNIST dataset?

8.2. Networks Using Blocks (VGG)

- VGG is based on simple, narrow, but deep network structure!
- VGG uses 3 x 3 filters, which is relatively small, so that we can extracts features in detail!

8.6. Residual Networks (ResNet) and ResNeXt

- It is not guaranteed that whenever we make CNN deeper, its perfomance is enhanced because of non-nested function classes!
- To solve this problem, we can introduce residual value (g(x) = f(x) x)!